

Imitation-learning-project

General Outline

The project aims to solve/build a single research idea by balancing its theory with empirical evaluation. We hope to begin by gaining intuition about the problem and addressing it on a simple toy task. The method can then be extended to non-trivial robot control tasks in order to compare its efficacy with baseline algorithms.

Multi-Agent Path Finding

- [Lifelong Multi-Agent Path Finding in A Dynamic Environment](#)
- Other papers by Reza

Deep Reinforcement Learning for Robot Control

- [C-Learning: Learning To Achieve Goals Via Recursive Classification](#)
- [COG: Connecting New Skills to Past Experience with Offline Reinforcement Learning](#)
- [Diversity Is All You Need: Learning Skills Without a Reward Function](#)
- [Hierarchical Deep Reinforcement Learning: Integrating Temporal Abstraction and Intrinsic Motivation](#)

Fundamental Reinforcement Learning

- [A Theoretical and Empirical Analysis of Expected Sarsa](#)
- [Expected Eligibility Traces](#)
- [Dyna: Model-based Planning Combined with Model-Free Learning](#)
- [Maxmin Q-Learning: Controlling The Estimation Bias Of Q-Learning](#)